

# problem4

November 8, 2023

## 1 EN.520.637 Foundations of Reinforcement Learning

Lab 4: Multi-armed Bandit and Monte Carlo Method (60 points)

### 1.1 Content

1. Multi-armed Bandit
2. Monte Carlo Method

```
[6]: %pylab inline
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.ticker as mtick
```

%pylab is deprecated, use %matplotlib inline and import the required libraries.  
Populating the interactive namespace from numpy and matplotlib

### 1.2 Problem Statement [P 21, Ch 2.3, Sutton]

Define a 10-armed bandit problem in which the action values  $q_*(a)$ ,  $a = 1, \dots, 10$ , are samples from a standard norm distribution, i.e. Gaussian distribution with mean = 0 and variance  $\sigma^2 = 1$ . Then, when selected  $A_t$  at time step  $t$ , the actual reward,  $R_t$  is selected from a Gaussian distribution with mean =  $q_*(A_t)$  and variance = 1.

### 1.3 Problem 1. Greedy and $\epsilon$ -greedy algorithm (15 points)

1. Implement a function/functions that run this game 2000 times with  $\epsilon$ -greedy algorithm. Your function/functions should take  $\epsilon$  as one of the inputs and output: (a) average reward at each time step (b) percentage of optimal action at each time step. (optimal action is defined by  $a^* = \underset{a}{\operatorname{argmax}} q^*(a)$ )
2. Call your function/functions to generate the average reward and percentage of optimal action at each time step with: (a) Greedy-algorithm (b)  $\epsilon$ -greedy algorithm,  $\epsilon = 0.01$  (c)  $\epsilon$ -greedy algorithm,  $\epsilon = 0.1$ .
3. Plot the average reward and percentage of optimal action of those three cases and compare with [P 23 Fig 2.2 Sutton].

### 1.3.1 1.

```
[ ]: def e_greedy(list_q, epsilon):
    opt_arm = np.argmax(list_q)
    avg_rewards, opt_actions = np.zeros(shape=(2000, 1000)), np.
    ↪ zeros(shape=(2000, 1000))

    #Repeat the game for 2000 independent runs
    for i in range(2000):
        total_opt_arm, total_reward = 0, 0
        list_Q, list_N = np.zeros(10), np.zeros(10)
        #Measure over 1000 steps
        for j in range(1000):
            #Exploitation
            if np.random.random() > epsilon:
                num_max = np.argwhere(list_Q == list_Q.max()).flatten()
                #if only one max
                if len(num_max) == 1:
                    tmp_arm = num_max[0]
                else:
                    tmp_arm = np.random.choice(num_max)
            #Exploration
            else:
                tmp_arm = np.random.randint(10)

            reward = np.random.normal(list_q[tmp_arm], 1, 1)[0]
            list_N[tmp_arm] += 1
            list_Q[tmp_arm] += (reward - list_Q[tmp_arm]) / list_N[tmp_arm]

            total_reward += reward
            if tmp_arm == opt_arm:
                total_opt_arm += 1
            avg_rewards[i][j] = reward
            # avg_rewards[i][j] = total_reward/(j+1)
            opt_actions[i][j] = total_opt_arm/(j+1)

        avg_rewards = avg_rewards.mean(axis = 0)
        opt_actions = opt_actions.mean(axis = 0)

    return avg_rewards, opt_actions
```

### 1.3.2 2.

```
[ ]: list_q = np.random.normal(0, 1, 10)

avg_rewards_a, opt_actions_a = e_greedy(list_q, 0)
avg_rewards_b, opt_actions_b = e_greedy(list_q, 0.01)
```


```
avg_rewards_c, opt_actions_c = e_greedy(list_q, 0.1)
```

### 1.3.3 3.


```
[ ]: steps = np.arange(1000)
plt.plot(steps, avg_rewards_a, label = "$\epsilon=0$ (greedy)", color = "green")
plt.plot(steps, avg_rewards_b, label = "$\epsilon=0.01$", color = "red")
plt.plot(steps, avg_rewards_c, label = "$\epsilon=0.1$", color = "blue")
plt.xlabel("Steps")
plt.ylabel("Average reward")
plt.title("The trend of this graph is similar to [P 23 Fig 2.2 Sutton] over 1000 time steps")
plt.legend()
plt.show()

plt.plot(steps, opt_actions_a, label = "$\epsilon=0$ (greedy)", color = "green")
plt.plot(steps, opt_actions_b, label = "$\epsilon=0.01$", color = "red")
plt.plot(steps, opt_actions_c, label = "$\epsilon=0.1$", color = "blue")
plt.xlabel("Steps")
plt.ylabel("% Optimal action")
plt.title("The trend of this graph is similar to [P 23 Fig 2.2 Sutton] over 1000 time steps")
plt.legend()
plt.show()
```

The trend of this graph is similar to [P 23 Fig 2.2 Sutton] over 1000 time steps



The trend of this graph is similar to [P 23 Fig 2.2 Sutton] over 1000 time steps



#### 1.4 Problem 2. UCB Action Selection (15 points)

1. Implement a function/functions that run this game 2000 times with UCB Action Selection algorithm. Your function/functions should take  $c$  as one of the inputs and output: - average reward at each time step.
2. Call your function/functions to generate the average reward at each time step with: - UCB Action Selection algorithm,  $c = 2$ .
3. Plot the average reward of 2.2 and 1.2c, then compare with [P 28 Fig 2.4 Sutton].

##### 1.4.1 1.

```
[ ]: def ucb(list_q, c):
    opt_arm = np.argmax(list_q)
    avg_rewards, opt_actions = np.zeros(shape=(2000, 1000)), np.
    ↪ zeros(shape=(2000, 1000))

    #Repeat the game for 2000 independent runs
    for i in range(2000):
        total_opt_arm, total_reward = 0, 0
        list_Q, list_N = np.zeros(10), np.zeros(10)
        #Measure over 1000 steps
```

```

for j in range(1000):
    zero_idx = np.where(list_N == 0)[0]
    if len(zero_idx) > 0:
        tmp_arm = zero_idx[0]
    else:
        tmp_arm = np.argmax(list_Q + c * np.sqrt(np.log(j+1)/list_N))

    reward = np.random.normal(list_q[tmp_arm], 1, 1)[0]
    list_N[tmp_arm] += 1
    list_Q[tmp_arm] += (reward - list_Q[tmp_arm]) / list_N[tmp_arm]

    total_reward += reward
    if tmp_arm == opt_arm:
        total_opt_arm += 1
    avg_rewards[i][j] = reward
    # avg_rewards[i][j] = total_reward/(j+1)
    opt_actions[i][j] = total_opt_arm/(j+1)

avg_rewards = avg_rewards.mean(axis = 0)
opt_actions = opt_actions.mean(axis = 0)

return avg_rewards, opt_actions

```

#### 1.4.2 2.

```
[ ]: avg_rewards_ucb, _ = ucb(list_q, 2)
```


#### 1.4.3 3.

```

[ ]: steps = np.arange(1000)
plt.plot(steps, avg_rewards_ucb, label = "UCB  $c=2$ ", color = "blue")
plt.plot(steps, avg_rewards_c, label = " $\epsilon$ -greedy  $\epsilon=0.1$ ",
    ↪ color = "gray")
plt.xlabel("Steps")
plt.ylabel("Average reward")
plt.title("The trend of this graph is similar to [P 28 Fig 2.4 Sutton] over
    ↪ 1000 time steps")
plt.legend()
plt.show()

```

The trend of this graph is similar to [P 28 Fig 2.4 Sutton] over 1000 time steps



### 1.5 3. Monte Carlo Method (CartPole-v1 environment) (30 points)

#### 1.5.1 3.1 CartPole Introduction

We now apply Monte Carlo Method for CartPole problem.

1. A pole is attached by an un-actuated joint to a cart, which moves along a frictionless track.
2. The system is controlled by applying a force of +1 or -1 to the cart.
3. The pendulum starts upright, and the goal is to prevent it from falling over.
4. A reward of +1 is provided for every timestep that the pole remains upright.
5. The episode ends when the pole is more than 15 degrees from vertical, or the cart moves more than 2.4 units from the center.
6. For more info (See [SOURCE ON GITHUB](#)).

The following examples show the basic usage of this testing environment:

#### 1.5.2 3.1.1 Episode initialization and Initial Value

```
[7]: import gym
```

```
[15]: env = gym.make('CartPole-v0')  
       observation = env.reset() ##Initial an episode
```

```

if gym.__version__>'0.26.0':
    observation = observation[0]

print("Initial observation is {}".format(observation))

print("\nThis means the cart current position is {}".format(observation[0]),
      end = '')
print(" with velocity {}".format(observation[1]))

print("and the pole current angular position is {}".format(observation[2]), end = '')
print(" with angular velocity {}".format(observation[3]))

```

Initial observation is [-0.03508221 -0.04424752 -0.02267548 0.03374215]

This means the cart current position is -0.03508221358060837 with velocity -0.044247522950172424,  
and the pole current angular position is -0.022675475105643272 with angular velocity 0.03374214842915535,

/home/gpu/.local/lib/python3.9/site-packages/gym/envs/registration.py:555:  
UserWarning: WARN: The environment CartPole-v0 is out of date. You should  
consider upgrading to version `v1`.  
logger.warn(

### 1.5.3 3.1.2 Take actions

Use env.step(action) to take an action

action is an integer from 0 to 1

0: “Left”; 1: “Right”

```

[16]: print("Current observation is {}".format(observation))

action = 0 #go left

##### simulate one step
if gym.__version__>'0.26.0':
    observation, reward, terminated, truncated, info = env.step(action)
    done = terminated or truncated
else:
    observation, reward, done, info = env.step(action)
#####

```

```

print("\nNew observation is {}".format(observation))
print("Step reward is {}".format(reward))
print("Did episode just ends? -{}".format(done)) # episode ends when 3.1(6)
↳ happens

```

Current observation is [-0.03508221 -0.04424752 -0.02267548 0.03374215]

New observation is [-0.03596716 -0.23903708 -0.02200063 0.31918535]

Step reward is 1.0

Did episode just ends? -False

### 1.5.4 3.1.3 Simulate multiple episodes

(You may uncomment those lines to see an animation. However, it will not work for JupyterHub since the animation requires GL instead of webGL. If you have Jupyter notebook locally on your computer, this version of code will work through a virtual frame.)

```

[17]: env = gym.make('CartPole-v0')
      observation = env.reset()
      total_reward = 0
      ep_num = 0
      # img = plt.imshow(env.render(mode='rgb_array'))

      for _ in range(1000):
          #     img.set_data(env.render(mode='rgb_array'))
          #     display.display(plt.gcf())
          #     display.clear_output(wait=True)

          action = env.action_space.sample() # this takes random actions

          ##### simulate one step
          if gym.__version__ > '0.26.0':
              observation, reward, terminated, truncated, info = env.step(action)
              done = terminated or truncated
          else:
              observation, reward, done, info = env.step(action)
          #####

          total_reward += reward

          if done:
              # episode just ends
              observation = env.reset() # reset episode
              ep_num += 1

      print("Average reward per episode is {}".format(total_reward/ep_num))

```



```
env.close()
```

Average reward per episode is 22.727272727272727

### 1.5.5 3.1.4 States Discretization

The class `DiscreteObs()` discretizes the observation space into discrete state space, based on `numpy.digitize` (Please read its description in <https://numpy.org/doc/stable/reference/generated/numpy.digitize.html>)

Discretization of observation space is necessary for tabular methods. You can use `DiscreteObs()` or any other library for discretizing the observation space.

```
[18]: class DiscretObs():

    def __init__(self, bins_list):
        self._bins_list = bins_list

        self._bins_num = len(bins_list)
        self._state_num_list = [len(bins)+1 for bins in bins_list]
        self._state_num_total = np.prod(self._state_num_list)

    def get_state_num_total(self):

        return self._state_num_total

    def obs2state(self, obs):

        if not len(obs)==self._bins_num:
            raise ValueError("observation must have length {}".format(self._
↪_bins_num))
        else:
            return [np.digitize(obs[i], bins=self._bins_list[i]) for i in
↪range(self._bins_num)]

    def obs2idx(self, obs):

        state = self.obs2state(obs)

        return self.state2idx(state)

    def state2idx(self, state):

        idx = 0
        for i in range(self._bins_num-1,-1,-1):
            idx = idx*self._state_num_list[i]+state[i]
```

```

    return idx

def idx2state(self, idx):

    state = [None]*self._bins_num
    state_num_cumul = np.cumprod(self._state_num_list)
    for i in range(self._bins_num-1,0,-1):
        state[i] = idx/state_num_cumul[i-1]
        idx -=state[i]*state_num_cumul[i-1]
    state[0] = idx%state_num_cumul[0]

    return state

# Recommended Discretization for Cartpole-v1 when using Monte-Carlo methods
bins_pos = np.linspace(-2.4,2.4,40)           # position
bins_d_pos = np.linspace(-3,3,5)              # velocity
bins_ang = np.linspace(-0.2618,0.2618,40)     # angle
bins_d_ang = np.linspace(-0.3,0.3,5)          # angular velocity

dobs = DiscretObs([bins_pos,bins_d_pos,bins_ang,bins_d_ang])
observation = env.reset()
if gym.__version__>'0.26.0':
    observation = observation[0]

state = dobs.obs2state(observation)
idx = dobs.obs2idx(observation)

print("Current position of the cart is {:.4f}\n".format(observation[0]))
print("Current velocity of the cart is {:.4f}\n".format(observation[1]))
print("Current angular position of the pole is {:.4f} rad\n".
      ↪format(observation[2]))
print("Current angular velocity of the pole is {:.4f} rad\n".
      ↪format(observation[3]))

print("which are mapped to state {}, with corresponding index {}".
      ↪format(state,idx))

```

Current position of the cart is 0.0490

Current velocity of the cart is 0.0004

Current angular position of the pole is -0.0015 rad

Current angular velocity of the pole is -0.0475 rad

which are mapped to state [20, 3, 20, 2], with corresponding index 25235

### 1.5.6 3.2 On-policy first-visit MC control

#### Task 3.2 [Coding, 30 points]

1. Implement “On-policy first-visit MC control” algorithm in [Ch 5.4 Sutton] to choose optimal actions
2. Simulate this algorithm for 40000 episodes.
3. Divide the previous 40000 episodes into 20 sets. Plot average rewards for each sets. (i.e. plot average rewards for the first 2000 episodes, the second 2000 episodes, ..., and the 15th 2000 episodes.)
4. Use greedy policy of the trained Q function to control the carpole for 100 episode, plot the accumulate rewards over 100 episode

```
[19]: ## Suggested flow (Feel free to modify and add)

# parameters for epsilon-greedy algorithm, when epsilon_decay_rate=1, the
↪algorithm implement a fixed
# epsilon value as epsilon_start, you can choose either fixed epsilon or
↪decaying epsilon

epsilon_start = 0.3
epsilon_decay_rate = 0.97

set_num = 20
s = 0
env = gym.make('CartPole-v1')

observation = env.reset()
if gym.__version__ > '0.26.0':
    observation = observation[0]

epsilon = epsilon_start    # set epsilon

ep_num = 0

#Initialize
policy = np.ones(shape = (dobs.get_state_num_total(), 2)) * 0.5
list_q = np.zeros(shape = (dobs.get_state_num_total(), 2))
list_returns = np.empty((list_q.shape[0], list_q.shape[1]), dtype=object)
list_returns.fill([])
result_mc = np.zeros(set_num)
total_reward = 0

episode = []
i = 0
while 1:
    current_state = dobs.obs2idx(observation)
```

```

action = np.random.choice(2, p=[policy[current_state][0],
↪policy[current_state][1]])

##### simulate one step
if gym.__version__>'0.26.0':
    observation, reward, terminated, truncated, info = env.step(action)
    done = terminated or truncated
else:
    observation, reward, done, info = env.step(action)
#####

episode.append((current_state, action, reward))
total_reward += reward
i += 1

if done: # end of episode
    all_state_action_pairs = [(s, a) for (s,a,r) in episode]
    G = 0
    for j in reversed(range(i)):
        tmp_state, tmp_action, tmp_reward = episode[j]
        G = 0.9 * G + tmp_reward
        if not (tmp_state, tmp_action) in all_state_action_pairs[0:j]:
            list_returns[tmp_state][tmp_action].append(G)
            list_q[tmp_state][tmp_action] = np.
↪mean(list_returns[tmp_state][tmp_action])
            opt_arm = np.argmax(list_q[tmp_state])
            for a in range(2):
                if a == opt_arm:
                    policy[tmp_state][a] = 1 - epsilon + (epsilon/2)
                else:
                    policy[tmp_state][a] = epsilon/2

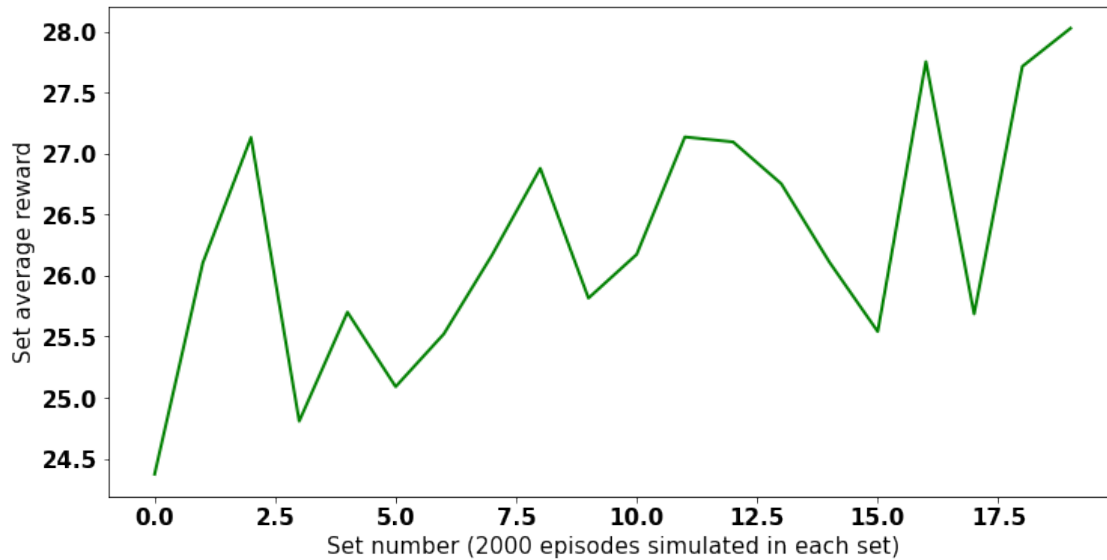
    observation = env.reset()
    if gym.__version__>'0.26.0':
        observation = observation[0]
    episode = []
    i = 0
    ep_num += 1
    if np.mod(ep_num,2000)==0: # end of every set of episode
        print(ep_num)
        epsilon = epsilon*epsilon_decay_rate # update epsilon
        result_mc[s] = total_reward/2000
        total_reward = 0
        s+=1
    if s == set_num:
        break

```

```
env.close()
```

```
2000  
4000  
6000  
8000  
10000  
12000  
14000  
16000  
18000  
20000  
22000  
24000  
26000  
28000  
30000  
32000  
34000  
36000  
38000  
40000
```

```
[20]: # check your result here (Feel free to modify)  
# the result_mc should be a (set_num, )-numpy array that records the average_  
# reward of a set of episodes  
  
# put your result here  
font = {'weight' : 'bold',  
        'size'   : 15}  
matplotlib.rc('font', **font)  
  
figure(figsize=(12,6))  
ax = subplot(1,1,1)  
ax.plot(range(0,set_num), result_mc, linewidth=2, color='g')  
  
plt.ylabel("Set average reward");  
plt.xlabel("Set number (2000 episodes simulated in each set)");
```



```
[23]: # Use greedy policy of the trained Q function to control the carpole for 100
      ↪ episode,
      # and plot the total reward received in each episode
      ## Suggested flow (Feel free to modify and add)

      env = gym.make('CartPole-v1')

      observation = env.reset()
      if gym.__version__ > '0.26.0':
          observation = observation[0]

      count = 0
      total_reward_mc = []
      total_reward = 0
      while 1:

          current_state = dobs.obs2idx(observation)    # discretize the observation
          ↪ space

          action = np.argmax(list_q[current_state])    # choose action by greedy
          ↪ policy of the trained Q

          ##### simulate one step
          if gym.__version__ > '0.26.0':
              observation, reward, terminated, truncated, info = env.step(action)
              done = terminated or truncated
          else:
              observation, reward, done, info = env.step(action)
```

```
#####

total_reward += reward

if done:

    observation = env.reset()
    if gym.__version__ > '0.26.0':
        observation = observation[0]

    count += 1

    total_reward_mc.append(total_reward)    # record the total reward until
    ↪ this episode
    # total_reward = 0
    if count == 100:
        break

font = {'weight' : 'bold',
        'size'   : 15}
matplotlib.rc('font', **font)

figure(figsize=(12,6))
ax = subplot(1,1,1)
ax.plot(range(100), total_reward_mc)

plt.ylabel("Accumulate Reward");
plt.xlabel("Episode");
```

