

# problem6

November 24, 2023

## 1 EN.520.637 Foundations of Reinforcement Learning

Lab 6: Deep Q-learning

Import torch, Gym and other necessary libraries

```
[10]: %pylab inline
import matplotlib.pyplot as plt

from torch import nn
import torch
import gym
from collections import deque
import itertools
import numpy as np
import random

env = gym.make('CartPole-v1')
```

Populating the interactive namespace from numpy and matplotlib

```
/usr/local/lib/python3.10/dist-packages/IPython/core/magics/pylab.py:159:
UserWarning: pylab import has clobbered these variables: ['random']
`%matplotlib` prevents importing * from pylab and numpy
  warn("pylab import has clobbered these variables: %s" % clobbered +
/usr/local/lib/python3.10/dist-packages/gym/core.py:317: DeprecationWarning:
WARN: Initializing wrapper in old step API which returns one bool instead
of two. It is recommended to set `new_step_api=True` to use new step API. This
will be the default behaviour in future.
  deprecation(
/usr/local/lib/python3.10/dist-
packages/gym/wrappers/step_api_compatibility.py:39: DeprecationWarning:
WARN: Initializing environment in old step API which returns one bool
instead of two. It is recommended to set `new_step_api=True` to use new step
API. This will be the default behaviour in future.
  deprecation(
```

```
[11]: GAMMA = 0.99
      BATCH_SIZE = 32
      BUFFER_SIZE = 50000
      MIN_REPLY_SIZE = 1000
      EPSILON_START = 1.0
      EPSILON_END = 0.02
      EPSILON_DECAY = 10000
      TARGET_UPDATE_FREQ = 1000
```

```
[12]: class Network(nn.Module):
      def __init__(self, env):
          super().__init__()

          in_features = int(np.prod(env.observation_space.shape))

          self.net = nn.Sequential(
              nn.Linear(in_features, 64),
              nn.Tanh(),
              nn.Linear(64, env.action_space.n))

      def forward(self, x):
          return self.net(x)

      def act(self, obs):
          obs_t = torch.as_tensor(obs, dtype=torch.float32)
          q_values = self(obs_t.unsqueeze(0))

          max_q_index = torch.argmax(q_values, dim=1)[0]
          action = max_q_index.detach().item()

          return action
```

```
[13]: def deep_Q_learning(Network):
      replay_buffer = deque(maxlen=BUFFER_SIZE)
      rew_buffer = deque([0.0], maxlen=100)
      set_size = 10
      step_list = deque([0.0], maxlen=set_size)

      episode = 0

      episode_reward = 0.0
      episode_step = 0.0

      n_set = 40
      result = np.zeros(n_set)
```

```

counter = 0

online_net = Network(env)
target_net = Network(env)

target_net.load_state_dict(online_net.state_dict())

optimizer = torch.optim.Adam(online_net.parameters(), lr=5e-4)

#initiallize replay buffer
obs = env.reset()
if gym.__version__>'0.26.0':
    obs = obs[0]
for _ in range(MIN_REPLY_SIZE):
    action = env.action_space.sample()

    ##### simulate one step
    if gym.__version__>'0.26.0':
        new_obs, rew, terminated, truncated, _ = env.step(action)
        done = terminated or truncated
    else:
        new_obs, rew, done, _ = env.step(action)
    #####

    transition = (obs, action, rew, done, new_obs)
    replay_buffer.append(transition)
    obs = new_obs

    if done:
        obs = env.reset()
        if gym.__version__>'0.26.0':
            obs = obs[0]

# Main Training Loop
obs = env.reset()
if gym.__version__>'0.26.0':
    obs = obs[0]

for step in itertools.count():
    epsilon = np.interp(step, [0, EPSILON_DECAY], [EPSILON_START, EPSILON_END])

    rnd_sample = random.random()

    if rnd_sample <= epsilon:

```

```

        action = env.action_space.sample()
    else:
        action = online_net.act(obs)

    ##### simulate one step
    if gym.__version__ > '0.26.0':
        new_obs, rew, terminated, truncated, _ = env.step(action)
        done = terminated or truncated
    else:
        new_obs, rew, done, _ = env.step(action)
    #####
    transition = (obs, action, rew, done, new_obs)
    replay_buffer.append(transition)
    obs = new_obs

    episode_reward += rew
    episode_step += 1

    if done:
        episode += 1

        obs = env.reset()
        if gym.__version__ > '0.26.0':
            obs = obs[0]
        step_list.append(episode_step)
        rew_buffer.append(episode_reward)
        episode_reward = 0.0
        episode_step = 0.0

    #Start Gradient Step
    transitions = random.sample(replay_buffer, BATCH_SIZE)

    obses = np.asarray([t[0] for t in transitions])
    actions = np.asarray([t[1] for t in transitions])
    rews = np.asarray([t[2] for t in transitions])
    dones = np.asarray([t[3] for t in transitions])
    new_obses = np.asarray([t[4] for t in transitions])

    obses_t = torch.as_tensor(obses, dtype=torch.float32)
    actions_t = torch.as_tensor(actions, dtype=torch.int64).unsqueeze(-1)
    rews_t = torch.as_tensor(rews, dtype=torch.float32).unsqueeze(-1)
    dones_t = torch.as_tensor(dones, dtype=torch.float32).unsqueeze(-1)
    new_obses_t = torch.as_tensor(new_obses, dtype=torch.float32)

    # Compute Targets
    target_q_values = target_net(new_obses_t)

```

```

max_target_q_values = target_q_values.max(dim=1, keepdim=True)[0]

targets = rews_t + GAMMA*(1 - dones_t)*max_target_q_values

# Compute Loss
q_values = online_net(obses_t)

action_q_values = torch.gather(input=q_values, dim=1, index=actions_t)

loss = nn.functional.smooth_l1_loss(action_q_values, targets)

#Gradient Descent
optimizer.zero_grad()
loss.backward()
optimizer.step()

# Update Target Network
if step % TARGET_UPDATE_FREQ == 0:
    target_net.load_state_dict(online_net.state_dict())

#Logging
if episode % (set_size+1) == 0:
    episode += 1
    result[counter] = np.mean(step_list)
    print("Finishing set {}/{}...".
    ↪format(counter+1,n_set),end="\r",flush=True)
    if counter == n_set-1:
        break
    counter+=1

return result, set_size

```

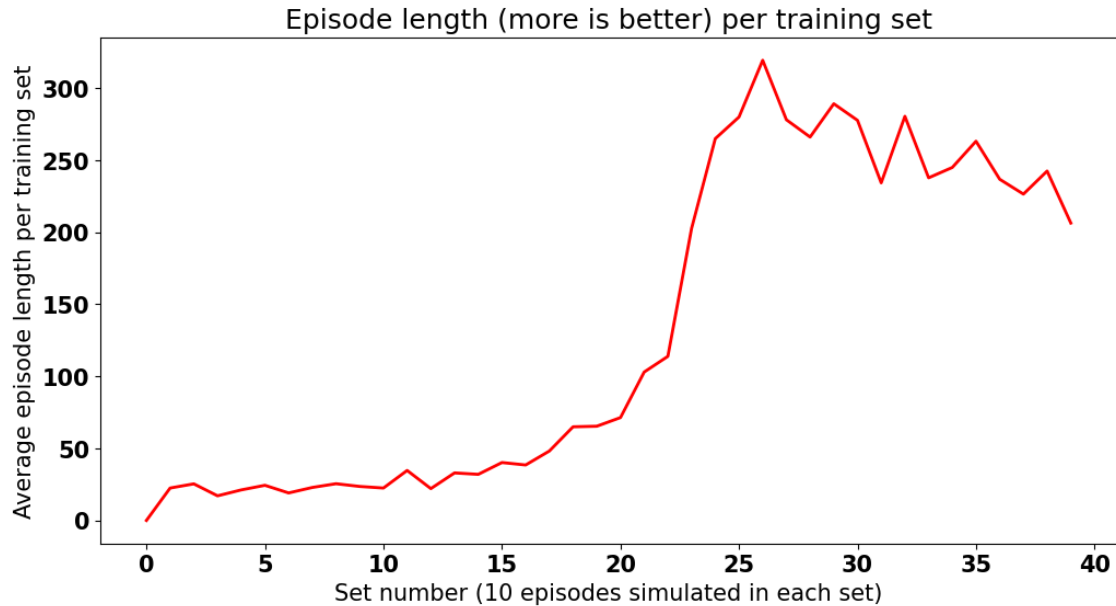
```

[14]: result, set_size = deep_Q_learning(Network)

font = {'weight' : 'bold',
        'size'    : 15}
matplotlib.rc('font', **font)

figure(figsize=(12,6))
ax = subplot(1,1,1)
ax.plot(range(len(result)), result, linewidth=2, color='r')
plt.title("Episode length (more is better) per training set")
plt.ylabel("Average episode length per training set");
plt.xlabel("Set number ({} episodes simulated in each set)".format(set_size));

```



0. Understand the Deep Q-learning codes above, a good reference can be found in Algorithm 1 Mnih, V., Kavukcuoglu, K., Silver, D. et al. Human-level control through deep reinforcement learning. Nature 518, 529–533 (2015)

1.0.1 [Task 1: 10 points]. Try to apply 2 different activation functions (e.g. Sigmoid, ReLU) in the Deep Q-learning codes above. Then, plot the average episode length (same plot as above).

```
[15]: class Network_Sigmoid(nn.Module):
    def __init__(self, env):
        super().__init__()

        in_features = int(np.prod(env.observation_space.shape))

        self.net = nn.Sequential(
            nn.Linear(in_features, 64),
            nn.Sigmoid(),
            nn.Linear(64, env.action_space.n))

    def forward(self, x):
        return self.net(x)

    def act(self, obs):
        obs_t = torch.as_tensor(obs, dtype=torch.float32)
        q_values = self(obs_t.unsqueeze(0))
```

```

        max_q_index = torch.argmax(q_values, dim=1)[0]
        action = max_q_index.detach().item()

    return action

class Network_ReLU(nn.Module):
    def __init__(self, env):
        super().__init__()

        in_features = int(np.prod(env.observation_space.shape))

        self.net = nn.Sequential(
            nn.Linear(in_features, 64),
            nn.ReLU(),
            nn.Linear(64, env.action_space.n))

    def forward(self, x):
        return self.net(x)

    def act(self, obs):
        obs_t = torch.as_tensor(obs, dtype=torch.float32)
        q_values = self(obs_t.unsqueeze(0))

        max_q_index = torch.argmax(q_values, dim=1)[0]
        action = max_q_index.detach().item()

    return action

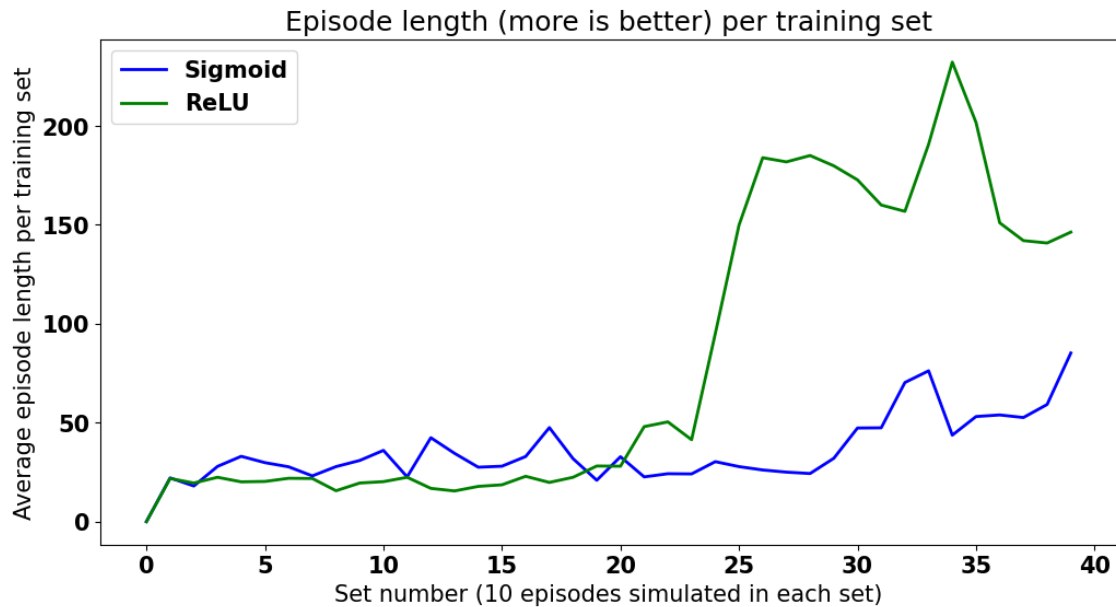
result_sigmoid, set_size = deep_Q_learning(Network_Sigmoid)
result_relu, set_size = deep_Q_learning(Network_ReLU)

font = {'weight' : 'bold',
        'size' : 15}
matplotlib.rc('font', **font)

figure(figsize=(12,6))
ax = subplot(1,1,1)
ax.plot(range(len(result_sigmoid)), result_sigmoid, linewidth=2, label = "Sigmoid", color='b')
ax.plot(range(len(result_relu)), result_relu, linewidth=2, label = "ReLU", color='g')
plt.title("Episode length (more is better) per training set")
plt.ylabel("Average episode length per training set");
plt.xlabel("Set number ({} episodes simulated in each set)".format(set_size));
plt.legend()

```

[15]: <matplotlib.legend.Legend at 0x7e784c07b4f0>



1.0.2 [Task 2: 10 points]. Using the best activation function you found in (1), try to modify the network structure defined in `self.net` to improve the average episode length. (e.g. add/decrease number of nodes/hidden layers) Plot the average episode length.

```
[19]: class Deep_Network_ReLU(nn.Module):
    def __init__(self, env):
        super().__init__()

        in_features = int(np.prod(env.observation_space.shape))

        self.net = nn.Sequential(
            nn.Linear(in_features, 128),
            nn.ReLU(),
            nn.Linear(128, env.action_space.n))

    def forward(self, x):
        return self.net(x)

    def act(self, obs):
        obs_t = torch.as_tensor(obs, dtype=torch.float32)
        q_values = self(obs_t.unsqueeze(0))

        max_q_index = torch.argmax(q_values, dim=1)[0]
        action = max_q_index.detach().item()
```



```

    return action

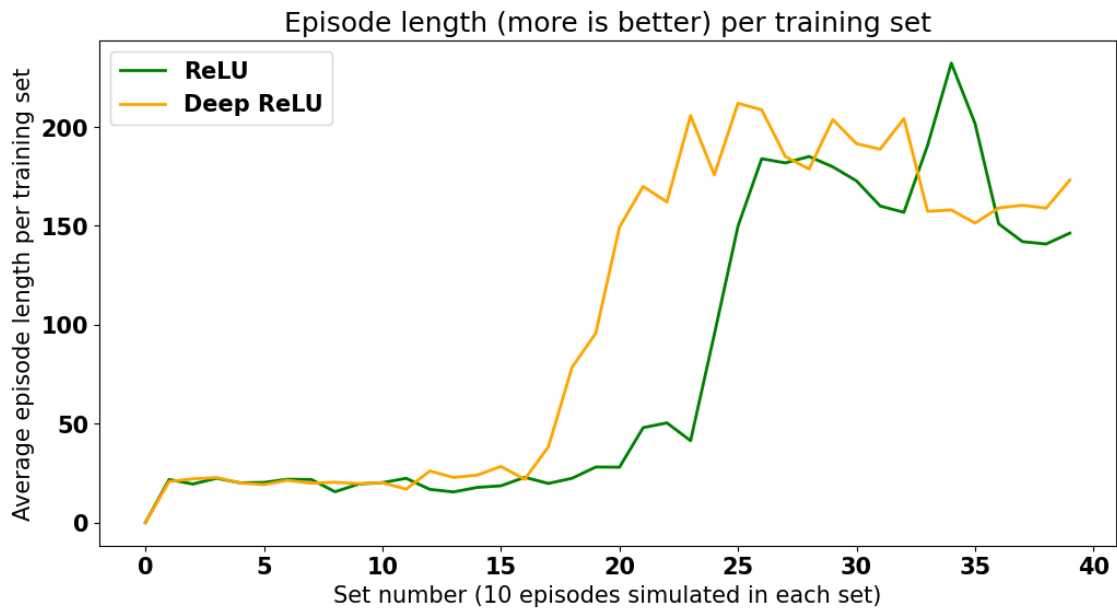
result_drelu, set_size = deep_Q_learning(Deep_Network_ReLU)

font = {'weight' : 'bold',
        'size'    : 15}
matplotlib.rc('font', **font)

figure(figsize=(12,6))
ax = subplot(1,1,1)
ax.plot(range(len(result_relu)), result_relu, linewidth=2, label = "ReLU",
        color='g')
ax.plot(range(len(result_drelu)), result_drelu, linewidth=2, label = "Deep
        ReLU", color='orange')
plt.title("Episode length (more is better) per training set")
plt.ylabel("Average episode length per training set");
plt.xlabel("Set number ({} episodes simulated in each set)".format(set_size));
plt.legend()

```

[19]: <matplotlib.legend.Legend at 0x7e784c0b8be0>



1.0.3 [Task 3: 10 points]. With the network structure and activation function unchanged, try to modify the reward to improve the average episode length. (e.g. give more reward if angle is small) Plot it.

```
[84]: def deep_Q_learning_3(Network):
    replay_buffer = deque(maxlen=BUFFER_SIZE)
    rew_buffer = deque([0.0],maxlen=100)
    set_size = 10
    step_list = deque([0.0],maxlen=set_size)

    episode = 0

    episode_reward = 0.0
    episode_step = 0.0

    n_set = 40
    result = np.zeros(n_set)
    counter = 0

    online_net = Network(env)
    target_net = Network(env)

    target_net.load_state_dict(online_net.state_dict())

    optimizer = torch.optim.Adam(online_net.parameters(), lr=5e-4)

    #initiallize replay buffer
    obs = env.reset()
    if gym.__version__>'0.26.0':
        obs = obs[0]
    for _ in range(MIN_REPLY_SIZE):
        action = env.action_space.sample()

        ##### simulate one step
        if gym.__version__>'0.26.0':
            new_obs, rew, terminated, truncated, _ = env.step(action)
            done = terminated or truncated
        else:
            new_obs, rew, done, _ = env.step(action)
        #####

        transition = (obs, action, rew, done, new_obs)
        replay_buffer.append(transition)
```

```

obs = new_obs

if done:
    obs = env.reset()
    if gym.__version__ > '0.26.0':
        obs = obs[0]

# Main Training Loop
obs = env.reset()
if gym.__version__ > '0.26.0':
    obs = obs[0]

for step in itertools.count():
    epsilon = np.interp(step, [0, EPSILON_DECAY], [EPSILON_START, EPSILON_END])

    rnd_sample = random.random()

    if rnd_sample <= epsilon:
        action = env.action_space.sample()
    else:
        action = online_net.act(obs)

    ##### simulate one step
    if gym.__version__ > '0.26.0':
        new_obs, rew, terminated, truncated, _ = env.step(action)
        done = terminated or truncated
    else:
        new_obs, rew, done, _ = env.step(action)
    #####
    transition = (obs, action, rew, done, new_obs)
    replay_buffer.append(transition)
    obs = new_obs

    episode_reward += rew
    episode_step += 1

    if done:
        episode += 1

        obs = env.reset()
        if gym.__version__ > '0.26.0':
            obs = obs[0]
        step_list.append(episode_step)
        rew_buffer.append(episode_reward)
        episode_reward = 0.0
        episode_step = 0.0

```

```

#Start Gradient Step
transitions = random.sample(replay_buffer, BATCH_SIZE)

obses = np.asarray([t[0] for t in transitions])
actions = np.asarray([t[1] for t in transitions])
rewards = np.asarray([t[2] for t in transitions])
dones = np.asarray([t[3] for t in transitions])
new_obses = np.asarray([t[4] for t in transitions])

#Modify reward by giving more reward if angle is small
angle = np.copy(new_obses[:, 2])
angle[angle > 0] = 0
rewards -= 0.5 * angle

obses_t = torch.as_tensor(obses, dtype=torch.float32)
actions_t = torch.as_tensor(actions, dtype=torch.int64).unsqueeze(-1)
rewards_t = torch.as_tensor(rewards, dtype=torch.float32).unsqueeze(-1)
dones_t = torch.as_tensor(dones, dtype=torch.float32).unsqueeze(-1)
new_obses_t = torch.as_tensor(new_obses, dtype=torch.float32)

# Compute Targets
target_q_values = target_net(new_obses_t)
max_target_q_values = target_q_values.max(dim=1, keepdim=True)[0]

targets = rewards_t + GAMMA*(1 - dones_t)*max_target_q_values

# Compute Loss
q_values = online_net(obses_t)

action_q_values = torch.gather(input=q_values, dim=1, index=actions_t)

loss = nn.functional.smooth_l1_loss(action_q_values, targets)

#Gradient Descent
optimizer.zero_grad()
loss.backward()
optimizer.step()

# Update Target Network
if step % TARGET_UPDATE_FREQ == 0:
    target_net.load_state_dict(online_net.state_dict())

#Logging
if episode % (set_size+1) == 0:
    episode += 1

```

```

        result[counter] = np.mean(step_list)
        print("Finishing set {}/{}...".
↪format(counter+1,n_set),end="\r",flush=True)
        if counter == n_set-1:
            break
        counter+=1

    return result, set_size

result_drelu_ir, set_size = deep_Q_learning_3(Deep_Network_ReLU)

font = {'weight' : 'bold',
        'size'    : 15}
matplotlib.rc('font', **font)

figure(figsize=(12,6))
ax = subplot(1,1,1)
ax.plot(range(len(result_drelu)), result_drelu, linewidth=2, label = "Deep_
↪ReLU", color='orange')
ax.plot(range(len(result_drelu_ir)), result_drelu_ir, linewidth=2, label = _
↪"Deep ReLU Improve Reward", color='purple')
plt.title("Episode length (more is better) per training set")
plt.ylabel("Average episode length per training set");
plt.xlabel("Set number ({} episodes simulated in each set)".format(set_size));
plt.legend()

```

[84]: <matplotlib.legend.Legend at 0x7e78466f3160>

